

Networks: IP and TCP

Internet Protocol

- Connectionless
 - Each packet is transported independently from other packets
- Unreliable
 - Delivery on a best effort basis
 - No acknowledgments
- Packets may be lost, reordered, corrupted, or duplicated
- IP packets
 - Encapsulate TCP and UDP packets
 - Encapsulated into link-layer frames

Data link frame


IP packet

TCP or UDP packet

IP Addresses and Packets

- IP addresses
 - IPv4: 32-bit addresses
 - IPv6: 128-bit addresses
- Address subdivided into **network**, **subnet**, and **host**
 - E.g., 128.148.32.110
- Broadcast addresses
 - E.g., 128.148.32.255
- Private networks
 - not routed outside of a LAN
 - 10.0.0.0/8
 - 172.16.0.0/12
 - 192.168.0.0/16

- IP header includes
 - Source address
 - Destination address
 - Packet length (up to 64KB)
 - Time to live (up to 255)
 - IP protocol version
 - Fragmentation information
 - Transport layer protocol information (e.g., TCP)



IP Address Space and ICANN

- Hosts on the internet must have unique IP addresses
- Internet Corporation for Assigned Names and Numbers
 - International nonprofit organization
 - Incorporated in the US
 - Allocates IP address space
 - Manages top-level domains
- Historical bias in favor of US corporations and nonprofit organizations
- Examples
 - 003/8 May 94 General Electric
 - 009/8 Aug 92 IBM
 - 012/8 Jun 95 AT&T Bell Labs
 - 013/8 Sep 91 Xerox Corporation
 - 015/8 Jul 94 Hewlett-Packard
 - 017/8 Jul 92 Apple Computer
 - 018/8 Jan 94 MIT
 - 019/8 May 95 Ford Motor
 - 040/8 Jun 94 Eli Lilly
 - 043/8 Jan 91 Japan Inet
 - 044/8 Jul 92 Amateur Radio Digital
 - 047/8 Jan 91 Bell-Northern Res.
 - 048/8 May 95 Prudential Securities
 - 054/8 Mar 92 Merck
 - 055/8 Apr 95 Boeing
 - 056/8 Jun 94 U.S. Postal Service

A Typical University's IP Space

- Most universities separate their network connecting dorms and the network connecting offices and academic buildings
- Dorms
 - Class B network 138.16.0.0/16 (64K addresses)
- Academic buildings and offices
 - Class B network 128.148.0.0/16 (64K addresses)
- CS department
 - Several class C (/24) networks, each with 254 addresses

IP Routing

- A router bridges two or more networks
 - Operates at the network layer
 - Maintains tables to forward packets to the appropriate network
 - Forwarding decisions based solely on the destination address
- Routing table
 - Maps ranges of addresses to LANs or other gateway routers


Internet Routes

- Internet Control Message Protocol (**ICMP**)
 - Used for network testing and debugging
 - Simple messages encapsulated in single IP packets
 - Considered a network layer protocol
- Tools based on ICMP
 - **Ping**: sends series of echo request messages and provides statistics on roundtrip times and packet loss
 - **Traceroute**: sends series ICMP packets with increasing TTL value to discover routes

ICMP Attacks

- Ping of death
 - ICMP specifies messages must fit a single IP packet (64KB)
 - Send a ping packet that exceeds maximum size using IP fragmentation
 - Reassembled packet caused several operating systems to crash due to a buffer overflow
- Smurf
 - Ping a broadcast address using a spoofed source address

Smurf Attack



IP Vulnerabilities


- Unencrypted transmission
 - Eavesdropping possible at any intermediate host during routing
- No source authentication
 - Sender can spoof source address, making it difficult to trace packet back to attacker
- No integrity checking
 - Entire packet, header and payload, can be modified while en route to destination, enabling content forgeries, redirections, and man-in-the-middle attacks
- No bandwidth constraints
 - Large number of packets can be injected into network to launch a denial-of-service attack
 - Broadcast addresses provide additional leverage

Denial of Service Attack

- Send large number of packets to host providing service
 - Slows down or crashes host
 - Often executed by botnet
- Attack propagation
 - Starts at zombies
 - Travels through tree of internet routers rooted
 - Ends at victim
- IP source spoofing
 - Hides attacker
 - Scatters return traffic from victim

Source:

M.T. Goodrich, [Probabilistic Packet Marking for Large-Scale IP Traceback](#), IEEE/ACM Transactions on Networking 16:1, 2008.



IP Traceback

- Problem
 - How to identify leaves of DoS propagation tree
 - Routers next to attacker
 - Issues
 - There are more than 2M internet routers
 - Attacker can spoof source address
 - Attacker knows that
 - Approaches
 - Filtering and tracing (immediate reaction)
 - Messaging (additional traffic)
 - Logging (additional storage)
 - Probabilistic marking
- traceback is being performed

Probabilistic Packet Marking

- Method
 - Random injection of information into packet header
 - Changes seldom used bits
 - Forward routing information to victim
 - Redundancy to survive packet losses
- Benefits
 - No additional traffic
 - No router storage
 - No packet size increase
 - Can be performed online or offline

Transmission Control Protocol

- TCP is a transport layer protocol guaranteeing reliable data transfer, in-order delivery of messages and the ability to distinguish data for multiple concurrent applications on the same host
- Most popular application protocols, including WWW, FTP and SSH are built on top of TCP
- TCP takes a stream of 8-bit byte data, packages it into appropriately sized segment and calls on IP to transmit these packets
- Delivery order is maintained by marking each packet with a **sequence number**
- Every time TCP receives a packet, it sends out an ACK to indicate successful receipt of the packet.
- TCP generally checks data transmitted by comparing a checksum of the data with a checksum encoded in the packet

Ports

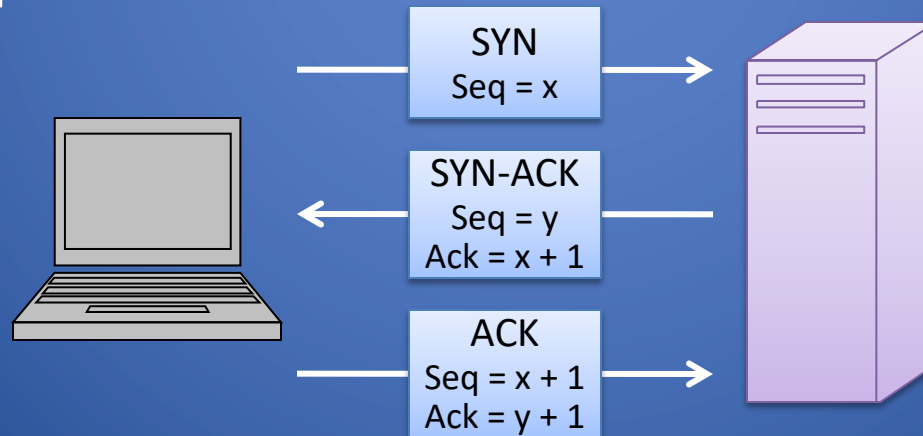
- TCP supports multiple concurrent applications on the same server
- Accomplishes this by having ports, 16 bit numbers identifying where data is directed
- The TCP header includes space for both a source and a destination port, thus allowing TCP to route all data
- In most cases, both TCP and UDP use the same port numbers for the same applications
- Ports 0 through 1023 are reserved for use by known protocols.
- Ports 1024 through 49151 are known as user ports, and should be used by most user programs for listening to connections and the like
- Ports 49152 through 65535 are private ports used for dynamic allocation by socket libraries

TCP Packet Format

Bit Offset	0-3	4-7	8-15	16-18	19-31
0	Source Port			Destination Port	
32	Sequence Number				
64	Acknowledgment Number				
96	Offset	Reserved	Flags	Window Size	
128	Checksum			Urgent Pointer	
160	Options				
>= 160	Payload				

Establishing TCP Connections

- TCP connections are established through a three way handshake.
- The server generally has a passive listener, waiting for a connection request
- The client requests a connection by sending out a SYN packet
- The server responds by sending a SYN/ACK packet, indicating an acknowledgment for the connection
- The client responds by sending an ACK to the server thus establishing connection



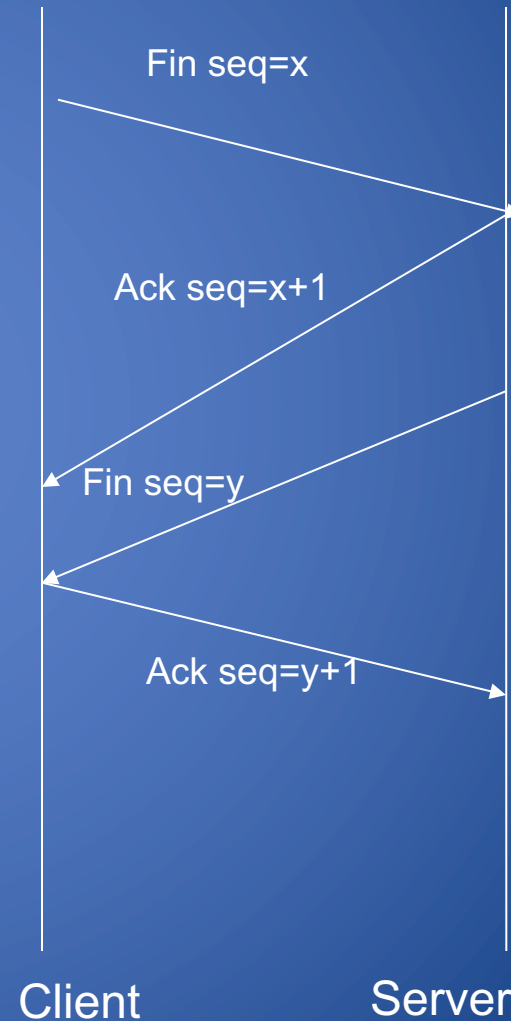
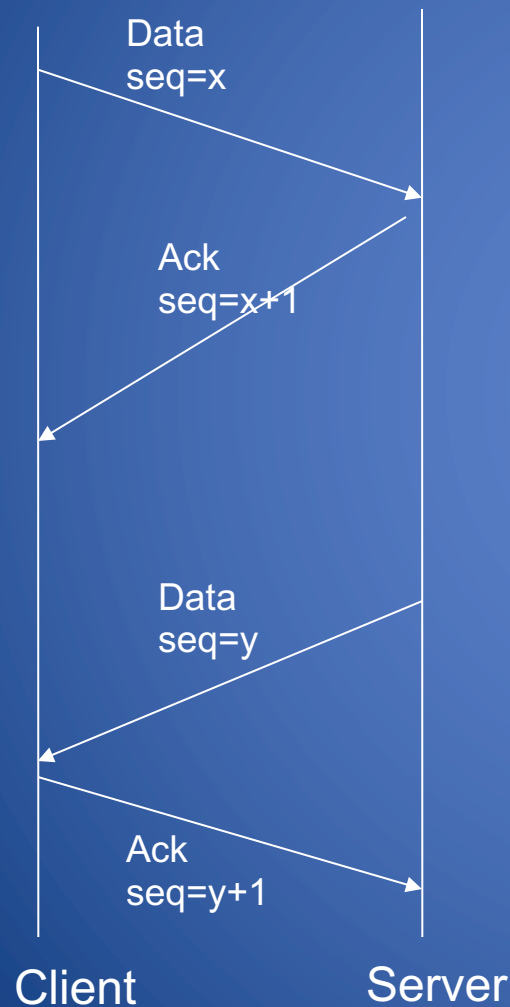
SYN Flood

- Typically DOS attack, though can be combined with other attack such as TCP hijacking
- Rely on sending TCP connection requests faster than the server can process them
- Attacker creates a large number of packets with spoofed source addresses and setting the SYN flag on these
- The server responds with a SYN/ACK for which it never gets a response (waits for about 3 minutes each)
- Eventually the server stops accepting connection requests, thus triggering a denial of service.
- Can be solved in multiple ways
- One of the common way to do this is to use SYN cookies

TCP Data Transfer

- During connection initialization using the three way handshake, initial sequence numbers are exchanged
- The TCP header includes a 16 bit checksum of the data and parts of the header, including the source and destination
- Acknowledgment or lack thereof is used by TCP to keep track of network congestion and control flow and such
- TCP connections are cleanly terminated with a 4-way handshake
 - The client which wishes to terminate the connection sends a FIN message to the other client
 - The other client responds by sending an ACK
 - The other client sends a FIN
 - The original client now sends an ACK, and the connection is terminated

TCP Data Transfer and Teardown



TCP Congestion Control

- During the mid-80s it was discovered that uncontrolled TCP messages were causing large scale network congestion
- TCP responded to congestion by retransmitting lost packets, thus making the problem worse
- What is predominantly used today is a system where ACKs are used to determine the maximum number of packets which should be sent out
- Most TCP congestion avoidance algorithms, avoid congestion by modifying a congestion window (cwnd) as more cumulative ACKs are received
- Lost packets are taken to be a sign of network congestion
- TCP begins with an extremely low cwnd and rapidly increases the value of this variable to reach bottleneck capacity
- At this point it shifts to a collision detection algorithm which slowly probes the network for additional bandwidth
- TCP congestion control is a good idea in general but allows for certain attacks.

Optimistic ACK Attack

- An optimistic ACK attack takes advantage of the TCP congestion control
- It begins with a client sending out ACKs for data segments it hasn't yet received
- This flood of optimistic ACKs makes the servers TCP stack believe that there is a large amount of bandwidth available and thus increase cwnd
- This leads to the attacker providing more optimistic ACKs, and eventually bandwidth use beyond what the server has available
- This can also be played out across multiple servers, with enough congestion that a certain section of the network is no longer reachable
- There are no practical solutions to this problem

Session Hijacking

- Also commonly known as TCP Session Hijacking
- A security attack over a protected network
- Attempt to take control of a network session
- Sessions are server keeping state of a client's connection
- Servers need to keep track of messages sent between client and the server and their respective actions
- Most networks follow the TCP/IP protocol
- IP Spoofing is one type of hijacking on large network

IP Spoofing

- IP Spoofing is an attempt by an intruder to send packets from one IP address that appear to originate at another
- If the server thinks it is receiving messages from the real source after authenticating a session, it could inadvertently behave maliciously
- There are two basic forms of IP Spoofing
 - Blind Spoofing
 - Attack from any source
 - Non-Blind Spoofing
 - Attack from the same subnet

Blind IP Spoofing

- The TCP/IP protocol requires that “acknowledgement” numbers be sent across sessions
- Makes sure that the client is getting the server’s packets and vice versa
- Need to have the right sequence of acknowledgment numbers to spoof an IP identity

Non-Blind IP Spoofing

- IP Spoofing without inherently knowing the acknowledgment sequence pattern
 - Done on the same subnet
 - Use a packet sniffer to analyze the sequence pattern
 - Packet sniffers intercept network packets
 - Eventually decodes and analyzes the packets sent across the network
 - Determine the acknowledgment sequence pattern from the packets
 - Send messages to server with actual client's IP address and with validly sequenced acknowledgment number

Packet Sniffers

- Packet sniffers “read” information traversing a network
 - Packet sniffers intercept network packets, possibly using ARP cache poisoning
 - Can be used as legitimate tools to analyze a network
 - Monitor network usage
 - Filter network traffic
 - Analyze network problems
 - Can also be used maliciously
 - Steal information (i.e. passwords, conversations, etc.)
 - Analyze network information to prepare an attack
- Packet sniffers can be either software or hardware based
 - Sniffers are dependent on network setup

Detecting Sniffers

- Sniffers are almost always passive
 - They simply collect data
 - They do not attempt “entry” to “steal” data
- This can make them extremely hard to detect
- Most detection methods require suspicion that sniffing is occurring
 - Then some sort of “ping” of the sniffer is necessary
 - It should be a broadcast that will cause a response only from a sniffer
- Another solution on switched hubs is ARP watch
 - An ARP watch monitors the ARP cache for duplicate entries of a machine
 - If such duplicates appear, raise an alarm
 - Problem: false alarms
 - Specifically, DHCP networks can have multiple entries for a single machine

Stopping Packet Sniffing

- The best way is to encrypt packets securely
 - Sniffers can capture the packets, but they are meaningless
 - Capturing a packet is useless if it just reads as garbage
 - SSH is also a much more secure method of connection
 - Private/Public key pairs makes sniffing virtually useless
 - On switched networks, almost all attacks will be via ARP spoofing
 - Add machines to a permanent store in the cache
 - This store cannot be modified via a broadcast reply
 - Thus, a sniffer cannot redirect an address to itself
- The best security is to not let them in in the first place
 - Sniffers need to be on your subnet in a switched hub in the first place
 - All sniffers need to somehow access root at some point to start themselves up

Port Knocking

- Broadly port knocking is the act of attempting to make connections to blocked ports in a certain order in an attempt to open a port
- Port knocking is fairly secure against brute force attacks since there are 65536^k combinations, where k is the number of ports knocked
- Port knocking however is very susceptible to replay attacks. Someone can theoretically record port knocking attempts and repeat those to get the same open port again
- One good way of protecting against replay attacks would be a time dependent knock sequence.

User Datagram Protocol

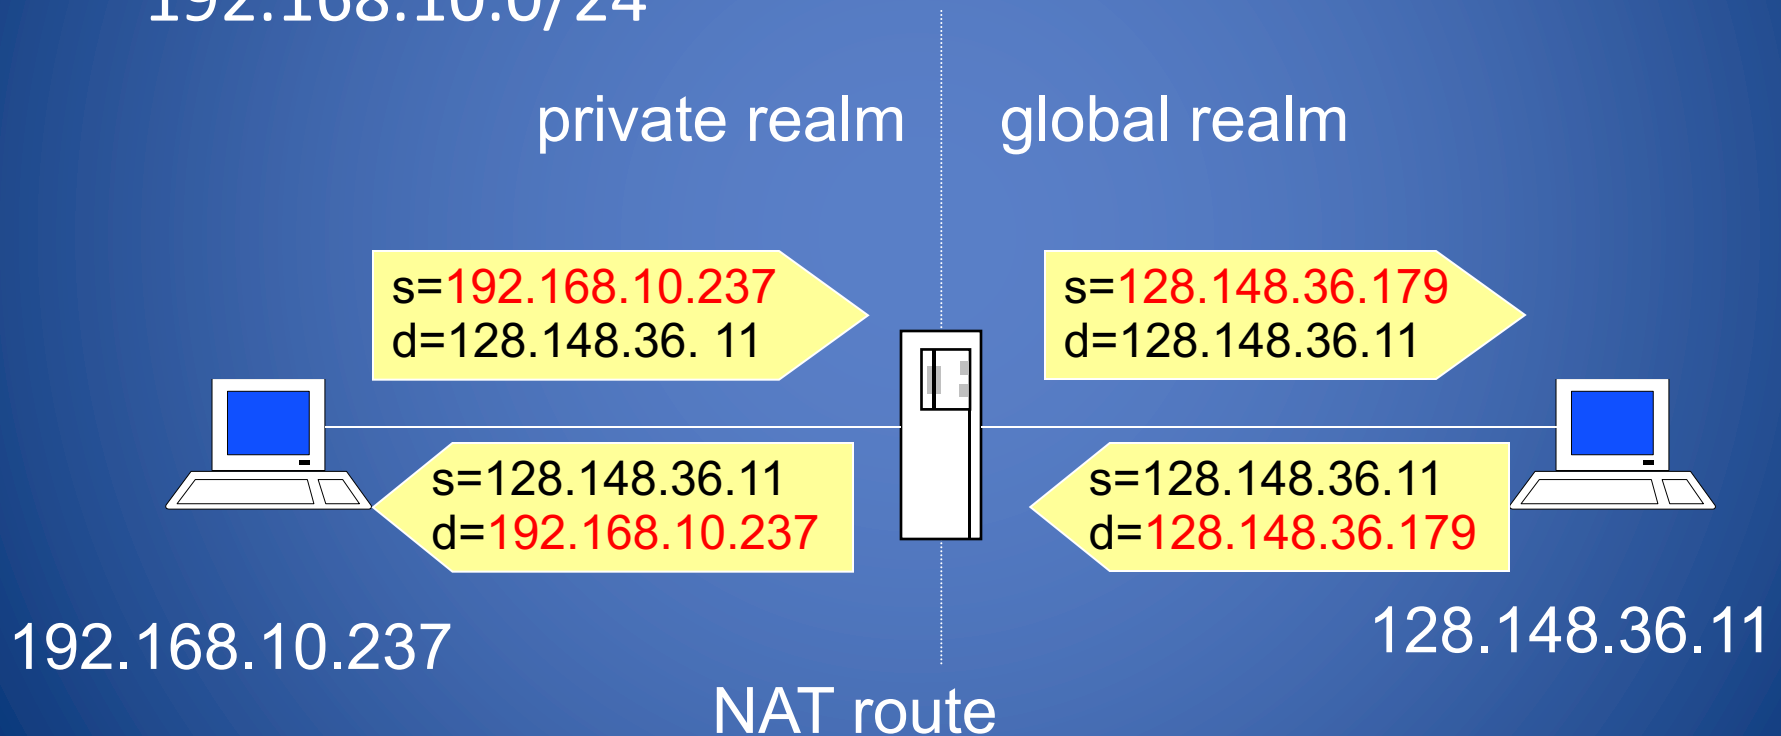
- UDP is a stateless, unreliable datagram protocol built on top of IP, that is it lies on level 4
- It does not provide delivery guarantees, or acknowledgments, but is significantly faster
- Can however distinguish data for multiple concurrent applications on a single host.
- A lack of reliability implies applications using UDP must be ready to accept a fair amount of error packages and data loss. Some application level protocols such as TFTP build reliability on top of UDP.
 - Most applications used on UDP will suffer if they have reliability. VoIP, Streaming Video and Streaming Audio all use UDP.
- UDP does not come with built in congestion protection, so while UDP does not suffer from the problems associated with optimistic ACK, there are cases where high rate UDP network access will cause congestion.

Network Address Translation

- Introduced in the early 90s to alleviate IPv4 address space congestion
- Relies on translating addresses in an internal network, to an external address that is used for communication to and from the outside world
- NAT is usually implemented by placing a router in between the internal private network and the public network.
- Saves IP address space since not every terminal needs a globally unique IP address, only an organizationally unique one
- While NAT should really be transparent to all high level services, this is sadly not true because a lot of high level communication uses things on IP

Translation

- Router has a pool of private addresses 192.168.10.0/24



IP Packet Modifications

